

Operating Systems Project 3 Travelogue

Noah Weiner

Part A

1. The different segments of MIPS memory are allocated to either kernel or user processes. Of the kernel segments, 2 are direct-mapped (cached or uncached) and one is TLB-mapped (cacheable).
2. (a) **tlb_random**: write a page to a random TLB slot.
(b) **tlb_write**: write to a specific TLB slot.
(c) **tlb_read**: read a page out of a slot.
(d) **tbl_probe**: find a TLB entry matching a virtual page number.
3. **PADDR_TO_KVADDR**: returns the physical address of a virtual address in kernel space (kseg0/kseg1).
4. The user stack pointer is initially set to USER-STACKTOP, which is MIPS_KSEG0, which is 0x80000000.
5. (a) **c0_entryhi**: the virtual page number of the current TLB entry, as well as its address space ID.
(b) **c0_entrylo**: the physical page number of the current TLB entry, as well as address space ID, as well as various flag bits (cacheable, dirty, global, valid).
6. The `as_*` functions are used to manipulate address spaces. The `as_prepare_load()` and `as_complete_load()` functions are called before and after loading an executable into memory, in order to set the correct page table permissions.
7. `vm_fault()` handles page faults.
8. (a) **0x100008**
 - i. Page no: 256 (0x100), offset: 8 (0x008)
 - ii. Translated address: 0x000008

iii. TLB contents:	Page #	Frame #
	0x100	0x008

(b) 0x101008

- i. Page no: 257 (0x101), offset: 8 (0x008)
- ii. Translated address: 0x001008

iii. TLB contents:	Page #	Frame #
	0x100	0x000
	0x101	0x001

(c) 0x1000f0

- i. Page no: 256 (0x100), offset: 240 (0x0f0)
- ii. Translated address: 0x0000f0

iii. TLB contents:	Page #	Frame #
	0x100	0x000
	0x101	0x001

(d) 0x41000

- i. Page no: 65 (0x41), offset: 0 (0x0)
- ii. Translated address: 0x002000

iii. TLB contents:	Page #	Frame #
	0x100	0x000
	0x101	0x001
	0x041	0x002

(e) 0x41b00

- i. Page no: 65 (0x41), offset: 2816 (0xb00)
- ii. Translated address: 0x002b00

iii. TLB contents:	Page #	Frame #
	0x100	0x000
	0x101	0x001
	0x041	0x002

Part B

1. MIPS_KSEG0 is 0x80000000
2. `firstpaddr` is the physical address of the first free page

3. `lastpaddr` is one past the end of the last free page
4. `firstfree` is the first free virtual address.
5. If there is 508 MB of RAM or more, the largest physical address is `0x1FC00000`.
6. `start.S` → `firstfree` → `ram_bootstrap` → `ramsize` → `lastpaddr` → `firstpaddr` → `ram_stealmem`
`ram_getsize` is never actually called.
7. Physical memory setup → polling hardware → VM initialization
8. `ram_bootstrap` → `alloc_kpages` → `ram_stealmem` → `vm_bootstrap`
 Again, I don't think `ram_getsize` is ever called.

Part C

1. `Dumbvm` has 12 pages of user stack, which it fills sequentially and then flops. It abuses functions (`as_prepare_load`, meant for preparing to read executables, and `ram_stealmem`, meant for vm bootstrapping) to do so.
2. `as_create` initializes the structure, `as_prepare_load` actually allocates memory via `getppages`, and memory is zeroed via `as_zero_region`.
3. It's not, yet, given that virtual memory isn't actually implemented. `PADDR_TO_KVADDR` goes the other way though.
4. Same as the above.
5. `kmalloc` grabs memory via `alloc_kpages` if necessary, otherwise `subpage_malloc`.
6. `kfree` tries `subpage_kfree`, and then `free_kpages` if necessary.
7. Each page maintains a list of subpages, which can be any size that's a multiple of 4 (preferably 8, according to `kmalloc.c`). Allocation is essentially finding a free subpage that fits whatever size is needed.
8. `as_create`, which allocates memory for the `adrspace` struct and essentially initializes all `adrspace` fields to 0.
9. `as_destroy`, which frees the `adrspace` struct (or would, if `kfree` worked).
10. To activate a user `adrspace`, `dumbvm` checks that the current address space is valid, and then invalidates all entries in the TLB.
11. A region is a set of pages designated for different access controls. There are two per address space.
12. `load_elf` uses `as_define_region` to create a separate region to load the executable into.
13. Nope, `adrspace` regions aren't protected in the `dumbvm`.
14. Regions CAN be accessed through the `adrspace` struct, but as far as I can most memory is accessed directly through the physical address.
15. Address space functions (`as_prepare_load`, `as_define_region`, etc) access `adrspace` regions for initialization.
16. `as_complete_load` currently does nothing. When project 3 is completed, it will likely set the `adrspace` region containing the executable bytecode to read-only, and any other post-elf-loading memory management to be done.
17. `as_define_stack` returns an initial pointer to the stack for a new process.
18. `runprogram()` in `runprogram.c`

Part D

1. `mips_trap` in `arch/mips/locore/trap.c`
2. Whenever a VM fault is triggered by trying to write read-only memory or a TLB miss.
3. It attempts to write the page into the TLB.

Summary Memory management is done via a two-level page hierarchy: each 4096 byte page maintains a list of sub-pages of varying length. Memory is divided into three segments, for user space and different types of kernel space. The kernel maintains a TLB which the MIPS uses for address translation. Processes own pages via an address space, which can have up to two regions for different access controls (for user processes, a read-only region for the executable byte code and a read-write address space). `Dumbvm` implements memory management as crudely as possible, and is meant to be replaced in project 3. It essentially abuses some functions meant for bootstrapping to allocate memory, and then crashes when physical memory is full. It doesn't handle de-allocation or regional access controls.

Collaborators and co-conspirators Eli, Zack, Kalen, Gavin, Nathaniel, Austin, Matthew, and Michelle